

# Introduction to Machine Learning

28 May 2019

Vineeth N Balasubramanian



आई आई टी हैदराबाद  
IIT Hyderabad

# A few (not so, yet) recent quotes

- “A breakthrough in machine learning would be worth ten Microsofts”  
(Bill Gates, Chairman, Microsoft)
- “Machine learning is the next Internet” (Tony Tether, Director, DARPA)
  - “Machine learning is today’s discontinuity” (Jerry Yang, ex-CEO, Yahoo)

# What is Machine Learning?




# What is Machine Learning?


- Making predictions or decisions from data
- “Programming computers to optimize a performance criterion using example data or past experience” (Ethem Alpaydin, Machine Learning, 2010)
- “A computer program is said to learn from experience  $E$  with respect to some class of tasks  $T$  and performance measure  $P$ , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ .” (Tom Mitchell, Machine Learning, 1997)
- “Learning general models from a data of particular examples”
- “Build a model that is *a good and useful approximation* to the data.”

# Today

## Traditional Programming



## Machine Learning



Source: Domingos


## Related Terms

Machine Learning, Data Mining, Knowledge Discovery,  
Artificial Intelligence, Statistical Learning, Pattern Recognition,  
Computational Learning



# When is Machine Learning Used?

- Human expertise does not exist
  - E.g. navigating on Mars
- Humans are unable to explain their expertise
  - E.g. speech recognition
- Solution changes in time
  - E.g. routing on a computer network
- Solution needs to be adapted to particular cases
  - E.g. user biometrics
- Data is cheap and abundant; knowledge is expensive and scarce



# Applications of Machine Learning

From: cheapsales@buystufffromme.com  
To: ang@cs.stanford.edu  
Subject: Buy now!

Deal of the week! Buy now!  
Rolex w4tchs - \$100  
Medicine (any kind) - \$50  
Also low cost M0rgages  
available.

Spam


From: Alfred Ng  
To: ang@cs.stanford.edu  
Subject: Christmas dates?

Hey Andrew,  
Was talking to Mom about plans  
for Xmas. When do you get off  
work. Meet Dec 22?  
Alf

Non-spam



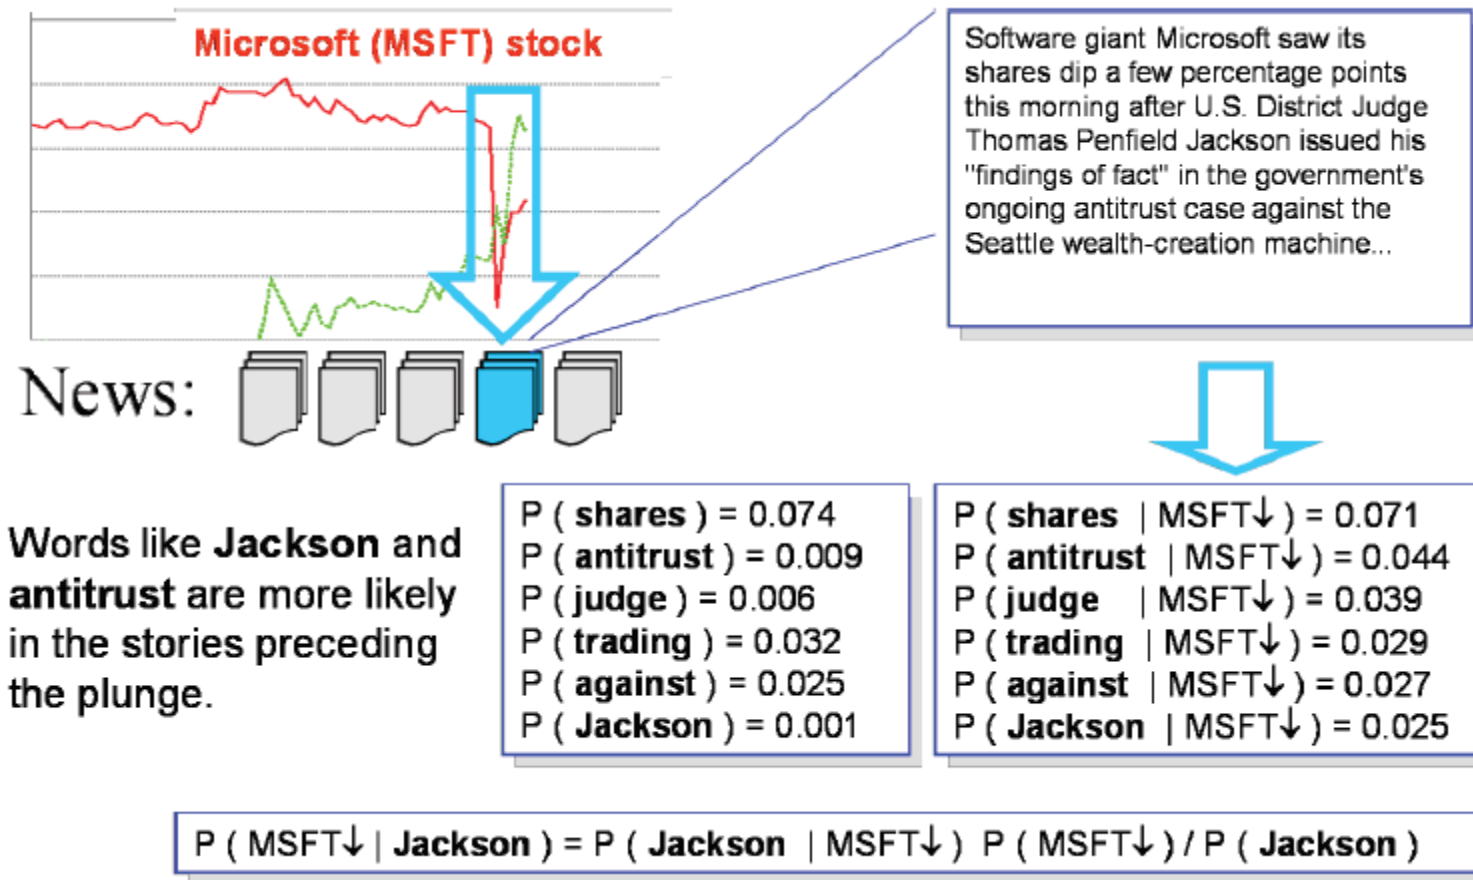
# Applications of Machine Learning



# Applications of Machine Learning

```
CCAGCTGCATCACAGGAGGCCAGCGAGCAGGTCTGTTCCAAGGGCCTTCGAGCCAGTCTG EI  
GAGGTGAAGGACGTCCTTCCCCAGGAGCCGGTGAGAAGCGCAGTCGGGGGCACGGGGATG EI  
TAAATTCTTCTGTTTGTTAACACCTTTCAGACTTATGTGTATGAAGGAGTAGAAGCCAAA IE  
AAACTAAAGAATTATTCTTTTACATTTTCAGTTTTTCTTGATCATGAAAACGCCAACAAAA IE  
AAAGCAGATCAGCTGTATAAACAGAAAATTATTCGTGGTTTCTGTCACTTGTGTGATGGT N  
TTGCCCTCAGCATCACCATGAACGGAGAGGCCATCGCCTGCGCTGAGGGCTGCCAGGCCA N
```

# Applications of Machine Learning



# Applications of Machine Learning

The screenshot displays the Netflix homepage with a red header. The main navigation bar includes links for 'Start Your 1 Month Free Trial', 'How It Works', 'Browse Selection', and '1 Month Free Trial Info'. A 'Member Sign In' button is located in the top right corner. The central banner features the text 'Unlimited TV episodes & movies instantly over the Internet plus unlimited DVDs by mail!'. Below this, three sections are highlighted: 'On your TV' (showing Wii, PS3, and XBOX 360 consoles), 'On your computer' (showing a laptop), and 'DVDs by mail' (showing a mailbox). Each section includes a 'Learn more >' button. A large 'PLUS' arrow points from the streaming options to the DVD by mail section. The bottom of the page is divided into three columns: 'FAQs' on the left, 'How does Netflix work?' in the center, and 'Start Your 1 Month Free Trial' on the right. The 'Start Your 1 Month Free Trial' section includes a form with fields for Email, Confirm Email, Password, and Confirm Password.

**NETFLIX** Member Sign In

**Start Your 1 Month Free Trial** **How It Works** **Browse Selection** **1 Month Free Trial Info**

Unlimited TV episodes & movies instantly over the Internet plus unlimited DVDs by mail!

**On your TV**  
Connect devices like these to your Netflix account to watch instantly on your TV.  
Wii PS3 XBOX 360  
Learn more > Learn more > Learn more >  
See other devices that stream instantly from Netflix

**On your computer**  
Watch instantly on your computer too!  
PLAY  
Learn more >

**DVDs by mail**  
Exchange as often as you want.  
No late fees - ever!

**FAQs**  
How does Netflix work?  
What is the selection like?  
How much does it cost?  
How many DVDs can I rent during my Free Trial?  
How fast will I get my DVDs?  
How long is the free trial?  
Can I cancel anytime?

**How does Netflix work?**  
**Rent what you want**  
Simply point and click to add movies & TV episodes to your list. Get DVDs by mail plus instantly watch movies (some new releases) & TV episodes (including current season) online on your PC or Mac or streamed instantly from Netflix over the Internet right to your TV via a Netflix ready device.

**Start Your 1 Month Free Trial**  
Free trial offer details.  
Email   
Confirm Email   
Password 4-10 characters   
Confirm Password

# More ML Applications

- Science (Astronomy, neuroscience, medical imaging, bio-informatics)
- Environment (energy, climate, weather, resources)
- Retail (Intelligent stock control, demographic store placement)
- Manufacturing (Intelligent control, automated monitoring, detection methods)
- Security (Intelligent smoke alarms, fraud detection)
- Marketing (promotions, ...)
- Management (Scheduling, timetabling)
- Finance (credit scoring, risk analysis...)
- Web data (information retrieval, information extraction, ...)

# More Recent ML Applications

- AlphaGo!
- Automating Employee Access Control
- Identifying whales in ocean based on audio recordings
- Predict wait times for patients in emergency rooms
- Extract heart failure diagnosis criteria from free-text physician notes
- Predicting hospital readmissions
- Is (s)he a psychopath?

Source: <http://www.forbes.com/sites/85broads/2014/01/06/six-novel-machine-learning-applications/#6b6f9a9e67bf>

# When are ML algorithms not needed?


- When the relationships between all system variables (input, output, and hidden) is completely understood!
- This is NOT the case for almost any real system!

# Overview of ML

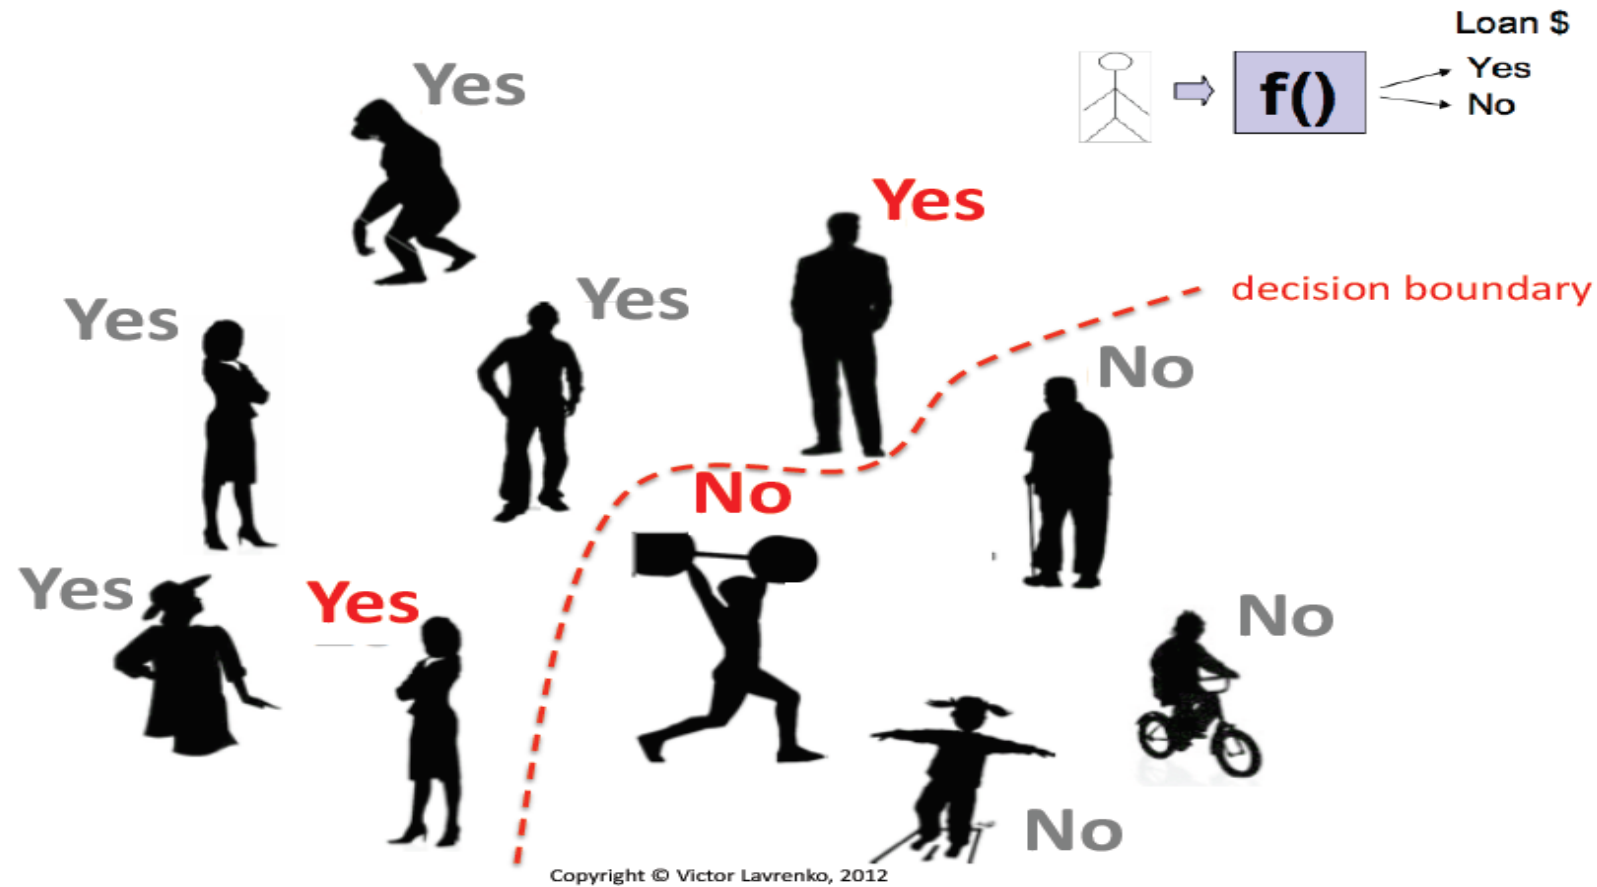
- Supervised learning
  - Predict an output  $y$  when given an input  $x$
  - For categorical  $y$  : classification.
  - For real-valued  $y$  : regression.
- Unsupervised learning
  - Create an internal representation of the input, e.g. clustering, dimensionality
  - This is important in machine learning as getting labels is often difficult and expensive
- Other settings of ML
  - Reinforcement learning (learning from “rewards”)
  - Semi-supervised learning (combines supervised + unsupervised)
  - Active learning, Transfer learning, Structured prediction




# Classification (Supervised Learning)



# Classification (Supervised Learning)




# Regression (Supervised Learning)




Copyright © Victor Lavrenko, 2012

# Clustering (Unsupervised Learning)



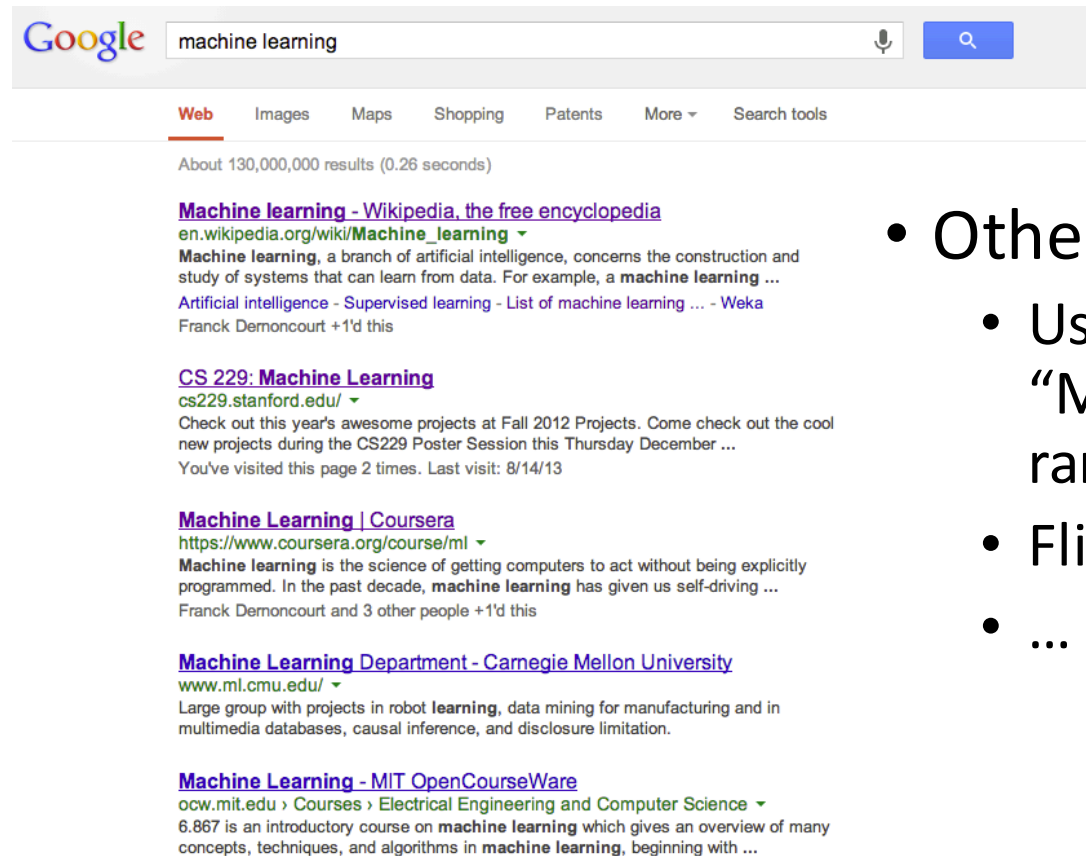
# Dimensionality Reduction (Unsupervised Learning)

- Large sample size is required for high-dimensional data
- Query accuracy and efficiency degrade rapidly as the dimension increases
- Strategies
  - Feature reduction
  - Feature selection
  - Manifold learning
  - Kernel learning



# Other Settings: Ranking (Supervised Learning)


Given a query and  
a set of web pages,  
rank them according  
to relevance



- Other applications
  - User preference, e.g. Netflix “My List” -- movie queue ranking
  - Flight search (search in general)
  - ...

# Other Settings: Reinforcement Learning

- Learning a policy: A **sequence** of outputs
- No supervised output but delayed reward
  - E.g. Game playing
  - E.g. Robot in a maze
- Multiple agents, partial observability, ...
- Example (Simple Demo):
  - <https://www.youtube.com/watch?v=DCjbk4m1G6I>



# ML Problems

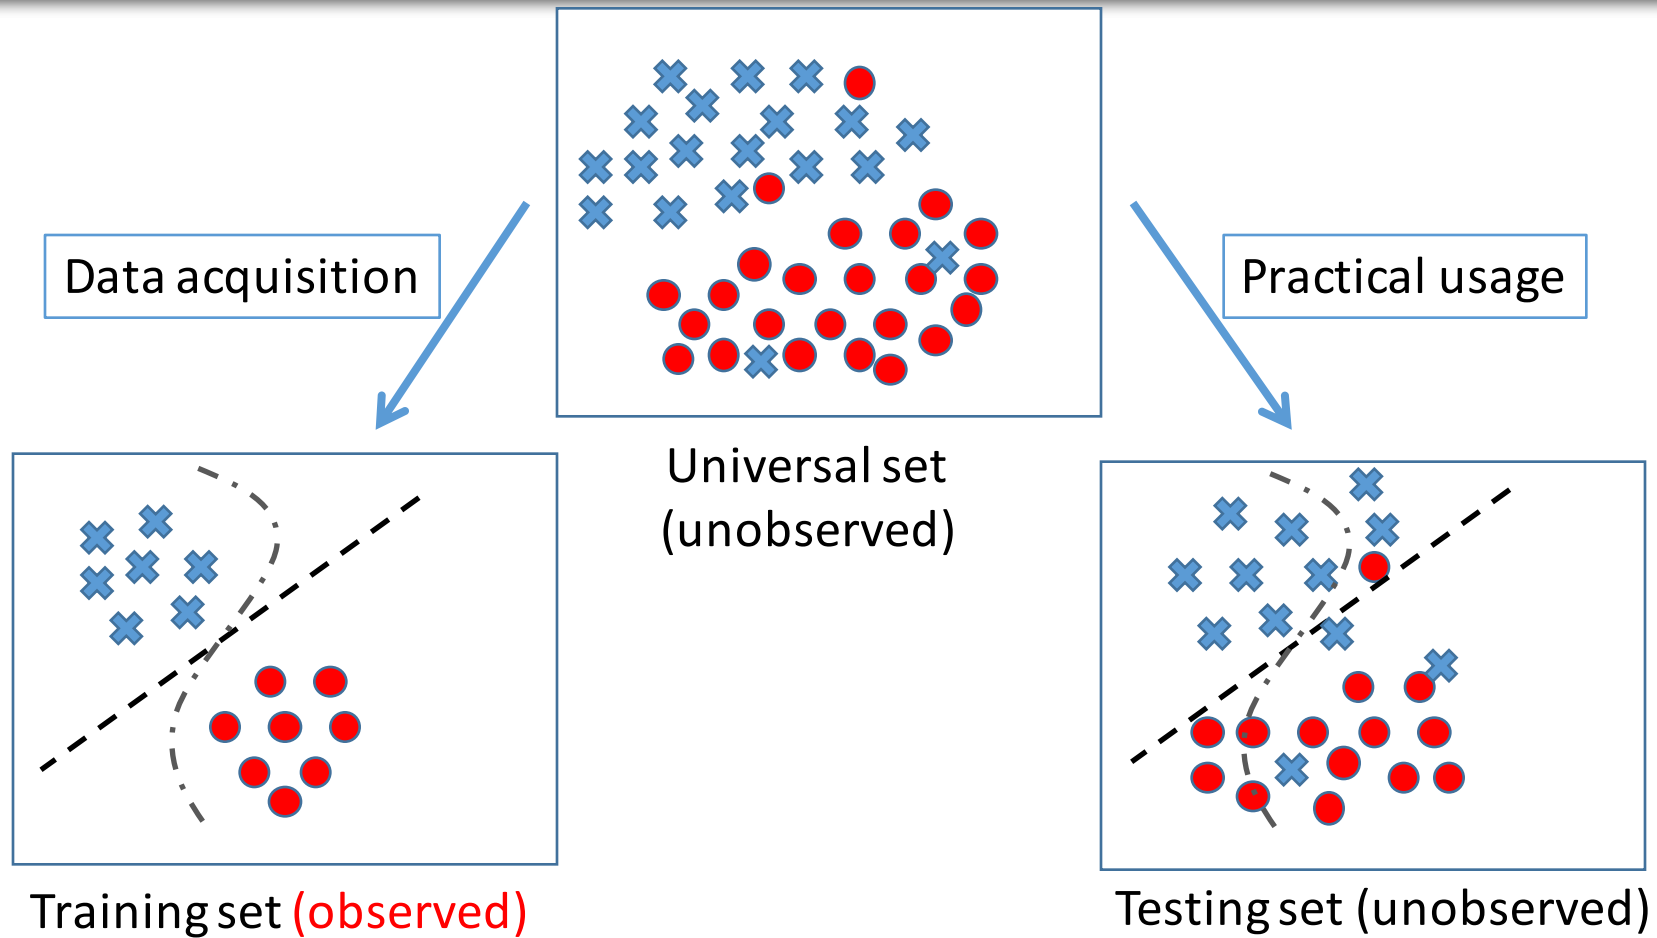
	<i>Supervised Learning</i>	<i>Unsupervised Learning</i>
<i>Discrete</i>	classification or categorization	clustering
<i>Continuous</i>	regression	dimensionality reduction



# ML in Practice


- Understanding domain, prior knowledge, and goals
- Data integration, selection, cleaning, pre-processing, etc.
- Learning models
- Interpreting results
- Consolidating and deploying discovered knowledge
- Loop

# Training and Testing ML Models



# Training and Testing ML Models

- Training is the process of making the system able to learn.
- **No free lunch** rule:
  - Training set and testing set may not come from the same distribution
  - Need to make some assumptions or bias



# Types of Models

- Inductive vs Transductive Learning
- Online vs Offline Learning
- Generative vs Discriminative Models
- Parametric vs Non-Parametric Models

# ML Datasets

- UCI Repository:  
<http://www.ics.uci.edu/~mlearn/MLRepository.html>
- Statlib: <http://lib.stat.cmu.edu/>
- Kaggle
- Many more...

# ML Resources

- MOOCs
  - Coursera, EdX, Udacity
- Conferences/Journals
  - JMLR, Machine Learning, IEEE Transactions on Neural Networks and Learning Systems, IEEE Transactions on Pattern Analysis and Machine Intelligence, Annals of Statistics
  - ICML, NIPS, KDD, IJCAI, AAAI, ICDM

# Mathematical Basis

- Functions, Logarithms and Exponentials
- Vectors, Dot Products, Orthogonality
- Matrices, Matrix Operations, Linear Transformations, Eigendecomposition
- Calculus, Differentiation, Integration
- Probability and Statistics
- Functional Analysis, Hilbert Spaces

# Programming

- Python
  - Numpy, Scipy – numerical/scientific computing, linear algebra
  - Matplotlib – for plotting
  - Scikitlearn – for machine learning



# Foundational Reading/Follow-up

- Math

- Part 1 of Deep Learning book: <http://www.deeplearningbook.org/>
- Essence of linear algebra: <http://youtu.be/kjBOesZCoqc>
- Essence of calculus: <https://goo.gl/Hnk1jA>

- Programming

- Practice Python
  - <https://try.jupyter.org/>
  - <https://docs.python.org/3/tutorial/>
  - Video Tutorials: <https://www.youtube.com/watch?v=cpPG0bKHYKc>
- Play with Numpy, Matplotlib, scikitlearn